

# Nengneng Yu

+1 484-995-8987 | [ynn1999@umd.edu](mailto:ynn1999@umd.edu) | College Park, MD

## PROFESSIONAL SUMMARY

- **Focused on data-driven system research:** leading projects that apply and implement data-driven approaches for computer systems, network, distributed systems, and biological data analysis. With hands-on experience on database systems and operating systems. Published and submitted works to conferences such as IEEE S&P and HotNet.
- **Competency:** Six year in computer science/engineering major in the US, fluent in English, and experienced in team collaboration.
- **Technical Proficiency:** C++/C, Java, Python, PyTorch, Pandas, Numpy, Linux, Machine Learning, Deep Learning, system programming, data structures, and algorithms.

## EDUCATION

- **University of Maryland College Park** College Park, MD  
*Doctor of Philosophy in Computer Science*  
*Advisor: Zaoxing(Alan) Liu <https://zaoxing.github.io/>*  
*GPA: 3.77/4.00*  
*Aug 2023 - Present*
- **Boston University** Boston, MA  
*Bachelor of Science in Computer Engineering*  
*Graduated with Magna Cum Laude*  
*Sep 2019 - May 2023*

## RESEARCH PROJECTS

- **Data-driven Generative Analysis System for Cancer Detection** Feb 2024 – Present  
*Mentor: Zaoxing(Alan) Liu, Yuefan Wang* *Froot Lab, UMD & Johns Hopkins Medicine*
  - Developed an AI-based generative system addressing the limitations of traditional bioinformatics in handling high-dimensional, small-sample proteomics data for gastric cancer detection and stage differentiation.
  - Designed a pipeline combining [Diffusion models](#) for synthetic data generation and [XGBoost](#) for biomarker identification, implemented using [PyTorch](#).
  - Achieved near 100% accuracy in gastric cancer NAT/Tumor classification. Improved CPTAC dataset performance by 10-20%, Korean dataset from 70% to over 90% on weighted accuracy among cancer stages.
- **Accelerating the mapping system for Micro Aerial Vehicle (MAV) navigation** Sep 2023 – Present  
*Mentor: Peiqing Chen* *UMD*
  - The performance of MAV autonomous navigation is limited by the slow updates and queries in mapping systems.
  - Developed a parallelized [software cache layer \(C++\)](#) to accelerate updates and queries in the mapping system.
  - Implemented the cached-mapping system on a [ROS-based MAV autonomous navigation simulation platform](#) and achieved up to 45% saving in end-to-end workflow runtime, raising flight speed by 26%.
- **Interactive Research Agents for Internet Incident Investigation** May 2023 – Nov 2023  
*Mentor: Zaoxing(Alan) Liu* *Froot Lab, UMD*
  - Developed an LLM-based agent to simulate experienced researchers and automate the investigation process, addressing the inefficiencies of traditional manual and time-consuming Internet incident investigations.
  - Built an agent using [Auto-GPT](#) and [GPT-4](#), equipped with autonomous [information retrieval](#), [knowledge memory](#), and [self-learning capabilities](#). Tested it on challenging scenarios such as the impact of hypothetical solar storms on networks.
  - Achieved 87.5% consistency in insights compared to human experts, effectively automating complex Internet incident analysis.
  - Co-first authored paper appeared at **HotNet 2023**
- **Advanced Persistent Threat (APT) Detection and Analysis** Feb 2022 – Present  
*Mentor: Zaoxing(Alan) Liu, Tuo Zhao* *Red Hat & Boston University & Georgia Tech*
  - Advanced Persistent Threats (APT) are stealthy and prolonged [cyber attacks](#) that target critical information or systems. Traditional detection methods struggle to identify new attack patterns and often lack precision due to data scarcity.
  - Built a flexible data pipeline with [Python](#) and [Pandas](#). Developed an APT detection system using a [Multi-Classification-Head Transformer](#) with a [pre-training](#) and [fine-tuning](#) workflow for efficient knowledge transfer.
  - Achieved 10%-29% improvement in fine-grained detection accuracy and 25%-50% reduction in training data needs across 8 real-world APT datasets. Enhanced model transparency, aiding in better post-attack investigation.
  - Submitted to **IEEE S&P 2025**

## COURSE PROJECTS

---

- **Concurrency Control Schemes for Database Systems** Jan 2023 – May 2024
  - Implemented and evaluated six concurrency control schemes, including [Two-Phase Locking \(2PL\)](#), [Optimistic Concurrency Control \(OCC\)](#), and [Multi-Version Concurrency Control \(MVCC\)](#). Developed two versions of 2PL with exclusive and shared locks, serial and parallel validation versions of OCC, and a simplified MVCC with [Serializable Snapshot Isolation \(SSI\)](#).
  - Developed and integrated the concurrency control schemes into a [main-memory key-value store](#) using [C++](#) and [thread management](#) techniques. Built a prototype transaction processing framework with a custom lock manager and multi-threaded execution support.
  - Conducted performance benchmarking with [CMake](#) and [CTest](#) to assess throughput and latency across varying transaction lengths and contention levels.
- **eBPF Modularity Project** Sep 2022 – Dec 2022
  - Collaborated with professors from Brown University and IBM engineers to advance research on building a comprehensive and reusable [eBPF module library](#).
  - Leveraged [OPENED](#) tool to analyze and decompose eBPF programs from open-source projects, extracting reusable modules for improved modularity and maintainability.
  - Designed a framework to transform extracted modules into a format compatible with [Bumblebee](#) tools for generating OCI images. Integrated “glue logic” for seamless module compatibility with [L3AF](#) and [Polycube](#).
- **Basic Unix-like Operating System** Jan 2022 – May 2022
  - Implemented a custom shell capable of executing commands via system calls, redirecting stdin/stdout, handling multiple command pipelines, and supporting background execution with “&”.
  - Developed a subset of the [POSIX threads API](#) in user mode, enabling multi-threaded execution with round-robin scheduling for effective resource management.
  - Designed and implemented a copy-on-write (COW) thread-local storage (TLS) mechanism to enable data sharing between threads while ensuring isolation of changes, improving memory management at the thread level.

## PUBLICATION & WORKS

---

- [1] Yajie Zhou\*, **Nengneng Yu\***, Zaoxing Liu, “Towards Interactive Simulacra of Internet Investigation by Human Researchers”, Hot Topics in Networks (HotNets), 2023
- [2] Yajie Zhou, **Nengneng Yu**, Simiao Zuo, Yue Yu, Haoming Yi, Chao Zhang, Tuo Zhao, Zaoxing Liu, “Fine-Grained, Adaptive Advanced Persistent Threats Detection with SENTINEL”, Under Submission

## SERVICES

---

- **University of Maryland College Park Department of Computer Science** Aug 2023 - Present  
*Graduate Teaching Assistant* *CMSC498B Cloud Computing*
- **University of Maryland College Park Department of Computer Science** Aug 2023 - Present  
*Graduate Teaching Assistant* *CMSC414 Network and Security*
- **University of Maryland College Park Department of Computer Science** Aug 2023 - Present  
*Graduate Teaching Assistant* *CMSC250 Discrete Structure*
- **Boston University College of Engineering** Sep 2022 - Dec 2022  
*Teaching Assistant* *EC440 Operating System*
- **Boston University College of Engineering** Jan 2022 - May 2022  
*Teaching Assistant* *EC414 Machine Learning*